



JIT-compiler in JVM seen by a Java developer

Vladimir Ivanov
HotSpot JVM Compiler
Oracle Corp.

MAKE THE
FUTURE
JAVA

ORACLE®

Agenda

- about compilers in general
 - ... and JIT-compilers in particular
- about JIT-compilers in HotSpot JVM
- monitoring JIT-compilers in HotSpot JVM

Static vs Dynamic

AOT vs JIT

Dynamic and Static Compilation Differences

- Static compilation
 - “ahead-of-time”(AOT) compilation
 - Source code → Native executable
 - Most of compilation work happens before executing
- Modern Java VMs use dynamic compilers (JIT)
 - “just-in-time” (JIT) compilation
 - Source code → Bytecode → Interpreter + JITted executable
 - Most of compilation work happens during executing

Dynamic and Static Compilation Differences

- Static compilation (AOT)
 - can utilize complex and heavy analyses and optimizations
 - ... but static information sometimes isn't enough
 - ... and it's hard to rely on profiling info, if any
 - moreover, how to utilize specific platform features (like SSE 4.2)?

Dynamic and Static Compilation Differences

- Modern Java VMs use dynamic compilers (JIT)
 - aggressive optimistic optimizations
 - through extensive usage of profiling info
 - ... but budget is limited and shared with an application
 - startup speed suffers
 - peak performance may suffer as well (not necessary)

JIT-compilation

- Just-In-Time compilation
- Compiled when needed
- Maybe immediately before execution
 - ...or when we decide it's important
 - ...or never?

Dynamic Compilation in JVM

JVM

- Runtime
 - class loading, bytecode verification, synchronization
- JIT
 - profiling, compilation plans, OSR
 - aggressive optimizations
- GC
 - different algorithms: throughput vs. response time

JVM: Makes Bytecodes Fast

- JVMs eventually JIT bytecodes
 - To make them fast
 - Some JITs are high quality optimizing compilers
- But cannot use existing static compilers directly:
 - Tracking OOPs (ptrs) for GC
 - Java Memory Model (volatile reordering & fences)
 - New code patterns to optimize
 - Time & resource constraints (CPU, memory)



JVM: Makes Bytecodes Fast

- JIT'ing requires Profiling
 - Because you don't want to JIT everything
- Profiling allows focused code-gen
- Profiling allows better code-gen
 - Inline what's hot
 - Loop unrolling, range-check elimination, etc
 - Branch prediction, spill-code-gen, scheduling



Dynamic Compilation (JIT)

- Knows about
 - loaded classes, methods the program has executed
- Makes optimization decisions based on code paths executed
 - Code generation depends on what is observed:
 - loaded classes, code paths executed, branches taken
- May re-optimize if assumption was wrong, or alternative code paths taken
 - Instruction path length may change between invocations of methods as a result of de-optimization / re-compilation

Dynamic Compilation (JIT)

- Can do non-conservative optimizations in dynamic
- Separates optimization from product delivery cycle
 - Update JVM, run the same application, realize improved performance!
 - Can be "tuned" to the target platform

Profiling

- Gathers data about code during execution
 - invariants
 - types, constants (e.g. null pointers)
 - statistics
 - branches, calls
- Gathered data is used during optimization
 - Educated guess
 - Guess can be wrong

Profile-guided optimization (PGO)

- Use profile for more efficient optimization
- PGO in JVMs
 - Always have it, turned on by default
 - Developers (usually) not interested or concerned about it
 - Profile is always consistent to execution scenario

Optimistic Compilers


- Assume profile is accurate
 - Aggressively optimize based on profile
 - Bail out if we're wrong
- ...and hope that we're usually right

Dynamic Compilation (JIT)

Overhead

- Is dynamic compilation overhead essential?
 - The longer your application runs, the less the overhead
- Trading off compilation time, not application time
 - Steal some cycles very early in execution
 - Done automagically and transparently to application
- Most of “perceived” overhead is compiler waiting for more data
 - ...thus running semi-optimal code for time being

JVM




Author: Alexey Shipilev

Mixed-Mode Execution

- Interpreted
 - Bytecode-walking
 - Artificial stack machine
- Compiled
 - Direct native operations
 - Native register machine

Bytecode Execution



Deoptimization

- Bail out of running native code
 - stop executing native (JIT-generated) code
 - start interpreting bytecode
- It's a complicated operation at runtime...

OSR: On-Stack Replacement

- Running method never exits?
- But it's getting really hot?
- Generally means loops, back-branching
- Compile and replace while running
- Not typically useful in large systems
- Looks great on benchmarks!



Optimizations

Optimizations in HotSpot JVM

- compiler tactics
 - delayed compilation
 - tiered compilation
 - on-stack replacement
 - delayed reoptimization
 - program dependence graph rep.
 - static single assignment rep.
- proof-based techniques
 - exact type inference
 - memory value inference
 - memory value tracking
 - constant folding
 - reassociation
 - operator strength reduction
 - null check elimination
 - type test strength reduction
 - type test elimination
 - algebraic simplification
 - common subexpression elimination
 - integer range typing
- flow-sensitive rewrites
 - conditional constant propagation
 - dominating test detection
 - flow-carried type narrowing
 - dead code elimination
- language-specific techniques
 - class hierarchy analysis
 - devirtualization
 - symbolic constant propagation
 - autobox elimination
 - escape analysis
 - lock elision
 - lock fusion
 - de-reflection
- speculative (profile-based) techniques
 - optimistic nullness assertions
 - optimistic type assertions
 - optimistic type strengthening
 - optimistic array length strengthening
 - untaken branch pruning
 - optimistic N-morphic inlining
 - branch frequency prediction
 - call frequency prediction
- memory and placement transformation
 - expression hoisting
 - expression sinking
 - redundant store elimination
 - adjacent store fusion
 - card-mark elimination
 - merge-point splitting
- loop transformations
 - loop unrolling
 - loop peeling
 - safepoint elimination
 - iteration range splitting
 - range check elimination
 - loop vectorization
- global code shaping
 - inlining (graph integration)
 - global code motion
 - heat-based code layout
 - switch balancing
 - throw inlining
- control flow graph transformation
 - local code scheduling
 - local code bundling
 - delay slot filling
 - graph-coloring register allocation
 - linear scan register allocation
 - live range splitting
 - copy coalescing
 - constant splitting
 - copy removal
 - address mode matching
 - instruction peepholing
 - DFA-based code generator

JVM: Makes Virtual Calls Fast

- C++ avoids virtual calls – because they are slow
- Java embraces them – and makes them fast
 - Well, mostly fast – JIT's do Class Hierarchy Analysis (CHA)
 - CHA turns most virtual calls into static calls
 - JVM detects new classes loaded, adjusts CHA
 - May need to re-JIT
 - When CHA fails to make the call static, inline caches
 - When IC's fail, virtual calls are back to being slow

Call Site

- The place where you make a call
- Monomorphic (“one shape”)
 - Single target class
- Bimorphic (“two shapes”)
- Polymorphic (“many shapes”)
- Megamorphic

Inlining

- Combine caller and callee into one unit
 - e.g. based on profile
 - ... or prove smth using CHA (Class Hierarchy Analysis)
 - Perhaps with a guard/test
- Optimize as a whole
 - More code means better visibility

Inlining

```
int addAll(int max) {  
    int accum = 0;  
    for (int i = 0; i < max; i++) {  
        accum = add(accum, i);  
    }  
    return accum;  
}  
  
int add(int a, int b) { return a + b; }
```

Inlining

```
int addAll(int max) {  
    int accum = 0;  
    for (int i = 0; i < max; i++) {  
        accum = accum + i;  
    }  
    return accum;  
}
```

Inlining and devirtualization

- Inlining is the most profitable compiler optimization
 - Rather straightforward to implement
 - Huge benefits: expands the scope for other optimizations
- OOP needs polymorphism, that implies virtual calls
 - Prevents naïve inlining
 - Devirtualization is required
 - (This does not mean you should not write OOP code)

JVM Devirtualization

- Developers shouldn't care
- Analyze hierarchy of currently loaded classes
- Efficiently devirtualize all monomorphic calls
- Able to devirtualize polymorphic calls
- JVM may inline dynamic methods
 - Reflection calls
 - Runtime-synthesized methods
 - JSR 292

Feedback multiplies optimizations

- On-line profiling and CHA produces information
 - ...which lets the JIT ignore unused paths
 - ...and helps the JIT sharpen types on hot paths
 - ...which allows calls to be devirtualized
 - ...allowing them to be inlined
 - ...expanding an ever-widening optimization horizon
- Result:

Large native methods containing tightly optimized machine code for hundreds of inlined calls!

Loop unrolling

```
public void foo(int[] arr, int a) {  
    for (int i = 0; i < arr.length; i++) {  
        arr[i] += a;  
    }  
}
```

Loop unrolling

```
public void foo(int[] arr, int a) {  
    for (int i = 0; i < arr.length; i=i+4) {  
        arr[i] += a; arr[i+1] += a; arr[i+2] += a; arr[i+3] += a;  
    }  
}
```

Loop unrolling

```
public void foo(int[] arr, int a) {  
    int new_limit = arr.length / 4;  
    for (int i = 0; i < new_limit; i++) {  
        arr[4*i] += a; arr[4*i+1] += a; arr[4*i+2] += a; arr[4*i+3] += a;  
    }  
}
```

```
for (int i = new_limit*4; i < arr.length; i++) {  
    arr[i] += a;  
}}
```

Lock Coarsening

```
public void m1(Object newValue) {  
    synchronized(this) {  
        field1 = newValue;  
    }  
    synchronized(this) {  
        field2 = newValue;  
    }  
}
```

Lock Coarsening

```
public void m1(Object newValue) {  
    synchronized(this) {  
        field1 = newValue;  
        field2 = newValue;  
    }  
}
```

Lock Eliding

```
public void m1() {  
    List list = new ArrayList();  
    synchronized (list) {  
        list.add(someMethod());  
    }  
}
```

Lock Eliding

```
public void m1() {  
    List list = new ArrayList();  
    synchronized (list) {  
        list.add(someMethod());  
    }  
}
```

Lock Eliding

```
public void m1() {  
    List list = new ArrayList();  
    list.add(someMethod());  
}
```


Escape Analysis

Initial version

```
public int m1() {  
    Pair p = new Pair(1, 2);  
    return m2(p);  
}  
  
public int m2(Pair p) {  
    return p.first + m3(p);  
}  
  
public int m3(Pair p) { return p.second;
```

Escape Analysis

After deep inlining

```
public int m1() {  
    Pair p = new Pair(1, 2);  
    return p.first + p.second;  
}
```

Escape Analysis

Optimized version

```
public int m1() {  
    return 3;  
}
```

Intrinsic

- Known to the JIT compiler
 - method bytecode is ignored
 - inserts “best” native code
- e.g. optimized sqrt in machine code
- Existing intrinsics
 - `String::equals`, `Math::*`, `System::arraycopy`, `Object::hashCode`, `Object::getClass`, `sun.misc.Unsafe::*`



HotSpot JVM

JVMs

- Oracle HotSpot
- IBM J9
- Oracle JRockit
- Azul Zing
- Excelsior JET
- Jikes RVM

HotSpot JVM

JIT-compilers

- client / C1
- server / C2
- tiered mode (C1 + C2)

HotSpot JVM

JIT-compilers

- client / C1
 - \$ java -client
 - only available in 32-bit VM
 - fast code generation of acceptable quality
 - basic optimizations
 - doesn't need profile
 - compilation threshold: 1,5k invocations

HotSpot JVM

JIT-compilers

- server / C2
 - \$ java -server
 - highly optimized code for speed
 - many aggressive optimizations which rely on profile
 - compilation threshold: 10k invocations

HotSpot JVM

JIT-compilers comparison

- Client / C1
 - + fast startup
 - peak performance suffers
- Server / C2
 - + very good code for hot methods
 - slow startup / warmup

Tiered compilation

C1 + C2

- -XX:+TieredCompilation
- Multiple tiers of interpretation, C1, and C2
- Level0=Interpreter
- Level1-3=C1
 - #1: C1 w/o profiling
 - #2: C1 w/ basic profiling
 - #3: C1 w/ full profiling
- Level4=C2

Monitoring JIT

Monitoring JIT-Compiler

- how to print info about compiled methods?
 - `-XX:+PrintCompilation`
- how to print info about inlining decisions
 - `-XX:+PrintInlining`
- how to control compilation policy?
 - `-XX:CompileCommand=...`
- how to print assembly code?
 - `-XX:+PrintAssembly`
 - `-XX:+PrintOptoAssembly` (C2-only)

Print Compilation

- -XX:+PrintCompilation
- Print methods as they are JIT-compiled
- Class + name + size

Print Compilation

Sample output

```
$ java -XX:+PrintCompilation
```

988	1	java.lang.String::hashCode (55 bytes)
1271	2	sun.nio.cs.UTF_8\$Encoder::encode (361 bytes)
1406	3	java.lang.String::charAt (29 bytes)

Print Compilation

Other useful info

- 2043 470 % ! jdk.nashorn.internal.ir.FunctionNode::accept @ 136 (265 bytes)
% == OSR compilation
! == has exception handles (may be expensive)
s == synchronized method
- 2028 466 n java.lang.Class::isArray (native)
n == native method

Print Compilation

Not just compilation notifications

- 621 160 java.lang.Object::equals (11 bytes) made not entrant
 - don't allow any new calls into this compiled version
- 1807 160 java.lang.Object::equals (11 bytes) made zombie
 - can safely throw away compiled version

No JIT At All?

- Code is too large
- Code isn't too «hot»
 - executed not too often

Print Inlining

- `-XX:+UnlockDiagnosticVMOptions -XX:+PrintInlining`
- Shows hierarchy of inlined methods
- Prints reason, if a method isn't inlined

Print Inlining

```
$ java -XX:+PrintCompilation -XX:+UnlockDiagnosticVMOptions -XX:+PrintInlining  
75  1      java.lang.String::hashCode (55 bytes)  
88  2      sun.nio.cs.UTF_8$Encoder::encode (361 bytes)  
      @ 14  java.lang.Math::min (11 bytes) (intrinsic)  
      @ 139 java.lang.Character::isSurrogate (18 bytes) never executed  
103 3      java.lang.String::charAt (29 bytes)
```

Inlining Tuning

- -XX:MaxInlineSize=35
 - Largest inlinable method (bytecode)
- -XX:InlineSmallCode=#
 - Largest inlinable compiled method
- -XX:FreqInlineSize=#
 - Largest frequently-called method...
- -XX:MaxInlineLevel=9
 - How deep does the rabbit hole go?
- -XX:MaxRecursiveInlineLevel=#
 - recursive inlining

Machine Code

- -XX:+PrintAssembly
- <http://wikis.sun.com/display/HotSpotInternals/PrintAssembly>
- Knowing code compiles is good
- Knowing code inlines is better
- Seeing the actual assembly is best!

-XX:CompileCommand=

- Syntax
 - “[command] [method] [signature]”
- Supported commands
 - **exclude** – never compile
 - **inline** – always inline
 - **dontinline** – never inline
- Method reference
 - class.name::methodName
- Method signature is optional



What Have We Learned?

- How JIT compilers work
- How HotSpot's JIT works
- How to monitor the JIT in HotSpot

Related Talks

“Quantum Performance Effects”

Sergey Kuksenko, Oracle


today, 13:30-14:30, «San-Francisco» hall

“Bulletproof Java Concurrency”

Aleksey Shipilev, Oracle

today, 15:30-16:30, «Moscow» hall

Questions?



vladimir.x.ivanov@oracle.com
@iwanowww

MAKE THE FUTURE JAVA



ORACLE®