

JVM for Dummies

(and for the rest of you, as well)

Intro

- Charles Oliver Nutter
 - “JRuby Guy”
 - Sun Microsystems 2006-2009
 - Engine Yard 2009-
- Primarily responsible for compiler, perf
 - Lots of bytecode generation
 - Lots of JIT monitoring

Today

- JVM Bytecode
 - Inspection
 - Generation
 - How it works
- JVM JIT
 - How it works
 - Monitoring
 - Assembly (don't be scared!)

Today

- JVM Bytecode
 - Inspection
 - Generation
 - How it works
- JVM JIT
 - How it works
 - Monitoring
 - Assembly



For Dummies

For people who want
to feel like Dummies

Part One: Bytecode

Bytecode Definition

- “... instruction sets designed for efficient execution by a software interpreter ...”
- “... suitable for further compilation into machine code.

Byte Code

- One-byte instructions
- 256 possible “opcodes”
- 200 in use on current JVMs
 - Room for more :-)
- Little variation since Java 1.0

Microsoft's CLR

- Stack-based, but not interpreted
- Two-byte “Wordcodes”
- Similar operations to JVM

Why Learn It

- Know your platform
 - Full understanding from top to bottom
- Bytecode generation is fun and easy
 - Build your own language?
- May need to read bytecode someday
 - Many libraries generate bytecode

Hello World

```
public class HelloWorld {  
    public static void main(String[] args) {  
        System.out.println("Hello, world");  
    }  
}
```

javap

- Java class file disassembler
- Basic operation shows class structure
 - Methods, superclasses, interface, etc
- -c flag includes bytecode
- -public, -private, -protected
- -verbose for stack size, locals, args

javap

```
~/projects/bytecode_for_dummies → javap HelloWorld  
Compiled from "HelloWorld.java"  
public class HelloWorld extends java.lang.Object{  
    public HelloWorld();  
    public static void main(java.lang.String[]);  
}
```

javap -c

~/projects/bytecode_for_dummies → javap -c HelloWorld

Compiled from "HelloWorld.java"

```
public class HelloWorld extends java.lang.Object{
```

```
public HelloWorld();
```

```
Code:
```

```
0:  aload_0
```

```
1:  invokespecial #1; //Method java/lang/Object."<init>":()V
```

```
4:  return
```

```
public static void main(java.lang.String[]);
```

```
Code:
```

```
0:  getstatic #2; //Field java/lang/System.out:Ljava/io/PrintStream;
```

```
3:  ldc #3; //String Hello, world
```

```
5:  invokevirtual #4; //Method java/io/PrintStream.println:  
                                     (Ljava/lang/String;)V
```

```
8:  return
```

```
}
```

javap -verbose

```
~/projects/bytecode_for_dummies → javap -c -verbose HelloWorld
Compiled from "HelloWorld.java"
public class HelloWorld extends java.lang.Object
  SourceFile: "HelloWorld.java"
  minor version: 0
  major version: 50
  Constant pool:
const #1 = Method      #6.#15; //  java/lang/Object."<init>":()V
const #2 = Field#16.#17; //  java/lang/System.out:Ljava/io/PrintStream;
const #3 = String      #18; //  Hello, world
const #4 = Method      #19.#20; //  java/io/PrintStream.println:(Ljava/lang/String;)V
const #5 = class#21; //  HelloWorld
...

{
```

javap -verbose

```
...
public HelloWorld();
  Code:
    Stack=1, Locals=1, Args_size=1
    0: aload_0
    1: invokespecial   #1; //Method java/lang/Object."<init>":()V
    4: return
 LineNumberTable:
    line 1: 0
```

javap -verbose

```
public static void main(java.lang.String[]);
```

```
Code:
```

```
Stack=2, Locals=1, Args_size=1
```

```
0:  getstatic #2; //Field java/lang/System.out:Ljava/io/PrintStream;
```

```
3:  ldc #3; //String Hello, world
```

```
5:  invokevirtual #4; //Method java/io/PrintStream.println:  
                                     (Ljava/lang/String;)V
```

```
8:  return
```

```
LineNumberTable:
```

```
line 3: 0
```

```
line 4: 8
```

```
}
```


Thank you!

Thank you!

(Just Kidding)

Let's try something a
little easier...

BiteScript

- (J)Ruby DSL for emitting JVM bytecode
 - Internal DSL
 - Primitive “macro” support
 - Reads like javap -c (but nicer)
- <http://github.com/headius/bitescript>

Installation

- Download JRuby from <http://jruby.org>
- Unpack, optionally add bin/ to PATH
 - Ahead of PATH if you have Ruby already
- [bin/]jruby -S gem install bitescript
- `bite myfile.bs` to run myfile.bs file
- `bitec myfile.bs` to compile myfile.bs file

BiteScript Users

- Mirah
 - Ruby-like language for writing Java code
 - BiteScript for JVM bytecode backend
- BrainF*ck implementation
- Other miscellaneous bytecode experiments

javap -c

~/projects/bytecode_for_dummies → javap -c HelloWorld

Compiled from "HelloWorld.java"

```
public class HelloWorld extends java.lang.Object{
```

```
public HelloWorld();
```

```
Code:
```

```
0:  aload_0
```

```
1:  invokespecial #1; //Method java/lang/Object."<init>":()V
```

```
4:  return
```

```
public static void main(java.lang.String[]);
```

```
Code:
```

```
0:  getstatic #2; //Field java/lang/System.out:Ljava/io/PrintStream;
```

```
3:  ldc #3; //String Hello, world
```

```
5:  invokevirtual #4; //Method java/io/PrintStream.println:  
                                     (Ljava/lang/String;)V
```

```
8:  return
```

```
}
```

BiteScript

```
main do
  getstatic java.lang.System, "out",
    java.io.PrintStream
  ldc "Hello, world!"
  invokevirtual java.io.PrintStream, "println",
    [java.lang.Void::TYPE, java.lang.Object]
  returnvoid
end
```


BiteScript

```
import java.lang.System
import java.io.PrintStream
```

JRuby's "import"
for Java classes

```
main do
  getstatic System, "out", PrintStream
  ldc "Hello, world!"
  invokevirtual PrintStream, "println", [void, object]
  returnvoid
end
```

Shortcuts for
void, int, string,
object, etc

BiteScript

```
main do  
  ldc "Hello, world!"  
  aprintln  
  returnvoid  
end
```



A BiteScript “macro”

BiteScript

```
macro :aprintln do
  getstatic System, "out", PrintStream
  swap
  invokevirtual PrintStream, "println",
    [void, object]
end
```

The Basics

- Stack machine
- Basic operations
- Flow control
- Class structures
- Exception handling

Stack Machine

- The “operand stack” holds operands
- Operations push and/or pop stack values
 - Exceptions: nop, wide, goto, jsr/ret
- Stack must be consistent
 - Largest part of bytecode verifier
- Stack is explicitly sized per method

The JVM Stack

```
import java.lang.System
import java.io.PrintStream

main do
  getstatic System, "out", PrintStream
  ldc "Hello, world!"
  invokevirtual PrintStream, "println",
    [void, object]
  returnvoid
end
```

| Depth | Value |
|-------|-------|
| 0 | |
| 1 | |
| 2 | |
| 3 | |
| 4 | |

The JVM Stack

```
import java.lang.System
import java.io.PrintStream

main do
  (getstatic System, "out", PrintStream)
  ldc "Hello, world!"
  invokevirtual PrintStream, "println",
    [void, object]
  returnvoid
end
```

| Depth | Value |
|-------|-------------------|
| 0 | out (a PS) |
| 1 | |
| 2 | |
| 3 | |
| 4 | |

The JVM Stack

```
import java.lang.System
import java.io.PrintStream

main do
  getstatic System, "out", PrintStream
  (ldc "Hello, world!")
  invokevirtual PrintStream, "println",
    [void, object]

  returnvoid
end
```

| Depth | Value |
|-------|------------------------|
| 0 | “Hello, world!” |
| 1 | out (a PS) |
| 2 | |
| 3 | |
| 4 | |

The JVM Stack

```
import java.lang.System
import java.io.PrintStream

main do
  getstatic System, "out", PrintStream
  ldc "Hello, world!"
  invokevirtual PrintStream, "println",
    [void, object]
  returnvoid
end
```

| Depth | Value |
|-------|-------|
| 0 | |
| 1 | |
| 2 | |
| 3 | |
| 4 | |

The JVM Stack

```
import java.lang.System
import java.io.PrintStream

main do
  getstatic System, "out", PrintStream
  ldc "Hello, world!"
  invokevirtual PrintStream, "println",
    [void, object]
  (returnvoid)
end
```

| Depth | Value |
|-------|-------|
| 0 | |
| 1 | |
| 2 | |
| 3 | |
| 4 | |

Basic Operations

- Stack manipulation
- Local variables
- Math
- Boolean

Stack Operations

| | | |
|------|---------|---|
| 0x00 | nop | Do nothing. |
| 0x57 | pop | Discard top value from stack |
| 0x58 | pop2 | Discard top two values |
| 0x59 | dup | Duplicate and push top value again |
| 0x5A | dup_x1 | Dup and push top value below second value |
| 0x5B | dup_x2 | Dup and push top value below third value |
| 0x5C | dup2 | Dup top two values and push |
| 0x5D | dup2_x1 | ...below second value |
| 0x5E | dup2_x2 | ...below third value |
| 0x5F | swap | Swap top two values |

Stack Juggling

dup
pop
swap
dup_x1
dup2_x2

| Depth | Value |
|-------|---------|
| 0 | value_0 |
| 1 | value_1 |
| 2 | |
| 3 | |
| 4 | |

Stack Juggling

dup
pop
swap
dup_x1
dup2_x2

| Depth | Value |
|-------|----------------|
| 0 | value_0 |
| 1 | value_0 |
| 2 | value_1 |
| 3 | |
| 4 | |

Stack Juggling

dup
pop
swap
dup_x1
dup2_x2

| Depth | Value |
|-------|---------|
| 0 | value_0 |
| 1 | value_1 |
| 2 | |
| 3 | |
| 4 | |

Stack Juggling

dup

pop

swap

dup_x1

dup2_x2

| Depth | Value |
|-------|----------------|
| 0 | value_1 |
| 1 | value_0 |
| 2 | |
| 3 | |
| 4 | |

Stack Juggling

dup

pop

swap

dup_x1

dup2_x2

| Depth | Value |
|-------|----------------|
| 0 | value_l |
| 1 | value_0 |
| 2 | value_l |
| 3 | |
| 4 | |

Stack Juggling

dup

pop

swap

dup_x1

dup2_x2

| Depth | Value |
|-------|----------------|
| 0 | value_l |
| 1 | value_0 |
| 2 | value_l |
| 3 | value_l |
| 4 | value_0 |

Typed Opcodes

<type><operation>

| | |
|---|-----------|
| b | byte |
| s | short |
| c | char |
| i | int |
| l | long |
| f | float |
| d | double |
| a | reference |

| |
|----------------------------------|
| Constant values |
| Local vars (load, store) |
| Array operations (aload, astore) |
| Math ops (add, sub, mul, div) |
| Boolean and bitwise |
| Comparisons |
| Conversions |

Where's boolean?

- Boolean is generally int 0 or 1
- Boolean operations push int 0 or 1
- Boolean branches expect 0 or nonzero
- To set a boolean...use int 0 or 1

Constant Values

| | | |
|-----------|----------------|--|
| 0x01 | aconst_null | Push null on stack |
| 0x02-0x08 | iload_[m1-5] | Push integer [-1 to 5] on stack |
| 0x09-0x0A | lconst_[0,1] | Push long [0 or 1] on stack |
| 0x0B-0x0D | fconst_[0,1,2] | Push float [0.0, 1.0, 2.0] on stack |
| 0x0E-0x0F | dconst_[0,1] | Push double [0.0, 1.0] on stack |
| 0x10 | bipush | Push byte value to stack as integer |
| 0x11 | sipush | Push short value to stack as integer |
| 0x12 | ldc | Push 32-bit constant to stack (int, float, string) |
| 0x14 | ldc2_w | Push 64-bit constant to stack (long, double) |

Why So Many?

- Reducing bytecode size
 - Special `iconst_0` and friends take no args
 - `bipush`, `sipush`: only 8, 16 bits arguments
- Pre-optimizing JVM
 - Specialized instructions can be optimized
 - Doesn't matter at all now

Constant Values

```
ldc "hello"  
dconst_1  
aconst_null  
bipush 4  
ldc_float 2.0
```

| Depth | Value |
|-------|-------|
| 0 | |
| 1 | |
| 2 | |
| 3 | |
| 4 | |
| 5 | |

Constant Values

ldc "hello"
dconst_1
aconst_null
bipush 4
ldc_float 2.0

| Depth | Value |
|-------|----------------|
| 0 | “hello” |
| 1 | |
| 2 | |
| 3 | |
| 4 | |
| 5 | |

Constant Values

```
ldc "hello"  
dconst_1  
aconst_null  
bipush 4  
ldc_float 2.0
```

| Depth | Value |
|-------|-------------|
| 0 | I.0d |
| 1 | |
| 2 | "hello" |
| 3 | |
| 4 | |
| 5 | |

Woah, Two Slots?

- JVM stack slots (and local vars) are 32-bit
- 64-bit values take up two slots
- “wide” before or “w” suffix
- 64-bit field updates not atomic!
 - Mind those concurrent longs/doubles!

Constant Values

```
ldc "hello"  
dconst_1  
aconst_null  
bipush 4  
ldc_float 2.0
```

| Depth | Value |
|-------|-------------|
| 0 | null |
| 1 | 1.0d |
| 2 | |
| 3 | “hello” |
| 4 | |
| 5 | |

Constant Values

```
ldc "hello"  
dconst_1  
aconst_null  
bipush 4  
ldc_float 2.0
```

| Depth | Value |
|-------|---------|
| 0 | 4 |
| 1 | null |
| 2 | 1.0d |
| 3 | |
| 4 | "hello" |
| 5 | |

Constant Values

ldc "hello"

dconst_1

aconst_null

bipush 4

ldc_float 2.0

| Depth | Value |
|-------|-------------|
| 0 | 2.0f |
| 1 | 4 |
| 2 | null |
| 3 | 1.0 |
| 4 | |
| 5 | "hello" |

Local Variable Table

- Local variables numbered from 0
 - Instance methods have “this” at 0
- Separate table maps numbers to names
- Explicitly sized in method definition

Local Variables

| | | |
|-----------|---------------|--|
| 0x15 | iload | Load integer from local variable onto stack |
| 0x16 | lload | ...long... |
| 0x17 | fload | ...float... |
| 0x18 | dload | ...double... |
| 0x19 | aload | ...reference... |
| 0x1A-0x2D | Packed loads | iload_0, aload_3, etc |
| 0x36 | istore | Store integer from stack into local variable |
| 0x37 | lstore | ...long... |
| 0x38 | fstore | ...float... |
| 0x39 | dstore | ...double... |
| 0x3A | astore | ...reference... |
| 0x3B-0x4E | Packed stores | fstore_2, dstore_0, etc |
| 0x84 | iinc | Add given amount to int local variable |

Local Variables

| Var | Value |
|-----|-------|
| 0 | |
| 1 | |
| 2 | |
| 3 | |
| 4 | |

```
ldc "hello"  
bipush 4  
istore 3  
dconst_0  
dstore 1  
astore 0  
aload 0  
iinc 3, 5
```

| Depth | Value |
|-------|-------|
| 0 | |
| 1 | |
| 2 | |
| 3 | |
| 4 | |

Local Variables

| Var | Value |
|-----|-------|
| 0 | |
| 1 | |
| 2 | |
| 3 | |
| 4 | |

ldc "hello"

bipush 4

istore 3

dconst_0

dstore 1

astore 0

aload 0

inc 3, 5

| Depth | Value |
|-------|----------------|
| 0 | "hello" |
| 1 | |
| 2 | |
| 3 | |
| 4 | |

Local Variables

| Var | Value |
|-----|-------|
| 0 | |
| 1 | |
| 2 | |
| 3 | |
| 4 | |

```
ldc "hello"  
bipush 4  
istore 3  
dconst_0  
dstore 1  
astore 0  
aload 0  
iinc 3, 5
```

| Depth | Value |
|-------|----------|
| 0 | 4 |
| 1 | "hello" |
| 2 | |
| 3 | |
| 4 | |

Local Variables

| Var | Value |
|-----|----------|
| 0 | |
| 1 | |
| 2 | |
| 3 | 4 |
| 4 | |

```
ldc "hello"  
bipush 4  
istore 3  
dconst_0  
dstore 1  
astore 0  
aload 0  
iinc 3, 5
```

| Depth | Value |
|-------|---------|
| 0 | "hello" |
| 1 | |
| 2 | |
| 3 | |
| 4 | |

Local Variables

| Var | Value |
|-----|-------|
| 0 | |
| 1 | |
| 2 | |
| 3 | 4 |
| 4 | |

```
ldc "hello"  
bipush 4  
istore 3  
dconst_0  
dstore 1  
astore 0  
aload 0  
iinc 3, 5
```

| Depth | Value |
|-------|------------|
| 0 | 0.0 |
| 1 | |
| 2 | "hello" |
| 3 | |
| 4 | |

Local Variables

| Var | Value |
|-----|------------|
| 0 | |
| 1 | 0.0 |
| 2 | |
| 3 | 4 |
| 4 | |

```
ldc "hello"  
bipush 4  
istore 3  
dconst_0  
dstore 1  
astore 0  
aload 0  
iinc 3, 5
```

| Depth | Value |
|-------|---------|
| 0 | "hello" |
| 1 | |
| 2 | |
| 3 | |
| 4 | |

Local Variables

| Var | Value |
|-----|----------------|
| 0 | “hello” |
| 1 | 0.0 |
| 2 | |
| 3 | 4 |
| 4 | |

```
ldc "hello"  
bipush 4  
istore 3  
dconst_0  
dstore 1  
astore 0  
aload 0  
iinc 3, 5
```

| Depth | Value |
|-------|-------|
| 0 | |
| 1 | |
| 2 | |
| 3 | |
| 4 | |

Local Variables

| Var | Value |
|-----|---------|
| 0 | "hello" |
| 1 | 0.0 |
| 2 | |
| 3 | 4 |
| 4 | |

```
ldc "hello"  
bipush 4  
istore 3  
dconst_0  
dstore 1  
astore 0  
aload 0  
iinc 3, 5
```

| Depth | Value |
|-------|----------------|
| 0 | "hello" |
| 1 | |
| 2 | |
| 3 | |
| 4 | |

Local Variables

| Var | Value |
|-----|---------|
| 0 | "hello" |
| 1 | 0.0 |
| 2 | |
| 3 | 9 |
| 4 | |

```
ldc "hello"  
bipush 4  
istore 3  
dconst_0  
dstore 1  
astore 0  
aload 0  
iinc 3, 5
```

| Depth | Value |
|-------|---------|
| 0 | "hello" |
| 1 | |
| 2 | |
| 3 | |
| 4 | |

Arrays

| | | |
|-----------|-------------------------|---|
| 0x2E-0x35 | [i,l,f,d,a,b,c,d]aload | Load [int, long, ...] from array (on stack) to stack |
| 0x4F-0x56 | [i,l,f,d,a,b,c,d]astore | Store [int, long, ...] from stack to array (on stack) |
| 0xBC | newarray | Construct new primitive array |
| 0xBD | anewarray | Construct new reference array |
| 0xBE | arraylength | Get array length |
| 0xC5 | multianewarray | Create multi-dimensional array |

Arrays

```
iconst_2  
newarray int  
dup  
iconst_0  
iconst_m1  
iastore  
iconst_0  
iaload
```

| Depth | Value |
|-------|-------|
| 0 | |
| 1 | |
| 2 | |
| 3 | |
| 4 | |
| 5 | |

Arrays

```
iconst_2  
newarray int  
dup  
iconst_0  
iconst_m1  
iastore  
iconst_0  
iaload
```

| Depth | Value |
|-------|----------|
| 0 | 2 |
| 1 | |
| 2 | |
| 3 | |
| 4 | |
| 5 | |

Arrays

```
iconst_2  
newarray int  
dup  
iconst_0  
iconst_m1  
iastore  
iconst_0  
iaload
```

| Depth | Value |
|-------|---------------------|
| 0 | int[2] {0,0} |
| 1 | |
| 2 | |
| 3 | |
| 4 | |
| 5 | |

Arrays

```
iconst_2  
newarray int  
dup  
iconst_0  
iconst_m1  
iastore  
iconst_0  
iaload
```

| Depth | Value |
|-------|---------------------|
| 0 | int[2] {0,0} |
| 1 | int[2] {0,0} |
| 2 | |
| 3 | |
| 4 | |
| 5 | |

Arrays

```
iconst_2  
newarray int  
dup  
(iconst_0)  
iconst_m1  
iastore  
iconst_0  
iaload
```

| Depth | Value |
|-------|--------------|
| 0 | 0 |
| 1 | int[2] {0,0} |
| 2 | int[2] {0,0} |
| 3 | |
| 4 | |
| 5 | |

Arrays

```
iconst_2  
newarray int  
dup  
iconst_0  
iconst_m1  
iastore  
iconst_0  
iaload
```

| Depth | Value |
|-------|--------------|
| 0 | -1 |
| 1 | 0 |
| 2 | int[2] {0,0} |
| 3 | int[2] {0,0} |
| 4 | |
| 5 | |

Arrays

```
iconst_2  
newarray int  
dup  
iconst_0  
iconst_m1  
(iastore)  
iconst_0  
iaload
```

| Depth | Value |
|-------|----------------|
| 0 | int[2] {-1, 0} |
| 1 | |
| 2 | |
| 3 | |
| 4 | |
| 5 | |

Arrays

```
iconst_2  
newarray int  
dup  
iconst_0  
iconst_m1  
iastore  
iconst_0  
iaload
```

| Depth | Value |
|-------|----------------|
| 0 | 0 |
| 1 | int[2] {-1, 0} |
| 2 | |
| 3 | |
| 4 | |
| 5 | |

Arrays

```
iconst_2  
newarray int  
dup  
iconst_0  
iconst_m1  
iastore  
iconst_0  
iaload
```

| Depth | Value |
|-------|-----------|
| 0 | -1 |
| 1 | |
| 2 | |
| 3 | |
| 4 | |
| 5 | |

Math Operations

| | add + | subtract - | multiply * | divide / | remainder % | negate -() |
|--------|----------|---------------|---------------|-------------|----------------|---------------|
| int | iadd | isub | imul | idiv | irem | ineg |
| long | ladd | lsub | lmul | ldiv | lrem | lneg |
| float | fadd | fsub | fmul | fdiv | frem | fneg |
| double | dadd | dsub | dmul | ddiv | drem | dneg |

Boolean and Bitwise

| | shift left | shift right | unsigned shift right | and | or | xor |
|-----|------------|-------------|-------------------------|------|-----|------|
| int | ishl | ishr | iushr | iand | ior | ixor |

Conversions

To:

From:

| | int | long | float | double | byte | char | short |
|--------|-----|------|-------|--------|------|------|-------|
| int | - | i2l | i2f | i2d | i2b | i2c | i2s |
| long | l2i | - | l2f | l2d | - | - | - |
| float | f2i | f2l | - | f2d | - | - | - |
| double | d2i | d2l | d2f | - | - | - | - |

Comparisons

| | | |
|------|-------|--|
| 0x94 | lcmp | Compare two longs, push int -1, 0, 1 |
| 0x95 | fcmpl | Compare two floats, push in -1, 0, 1 (-1 for NaN) |
| 0x96 | fcmpg | Compare two floats, push in -1, 0, 1 (1 for NaN) |
| 0x97 | dcmpl | Compare two doubles, push in -1, 0, 1 (-1 for NaN) |
| 0x98 | dcmpg | Compare two doubles, push in -1, 0, 1 (1 for NaN) |

Flow Control

- Inspect stack and branch
 - Or just branch, via goto
- Labels mark branch targets
- Wide variety of tests

Flow Control

| | | |
|------|-----------|---|
| 0x99 | ifeq | If zero on stack, branch |
| 0x9A | ifne | If nonzero on stack, branch |
| 0x9B | iflt | If stack value is less than zero, branch |
| 0x9C | ifge | If stack value is greater than or equal to zero, branch |
| 0x9D | ifgt | If stack value is greater than zero, branch |
| 0x9E | ifle | If stack value is less than or equal to zero, branch |
| 0x9F | if icmpeq | If two integers on stack are eq, branch |
| 0xA0 | if icmpne | If two integers on stack are ne, branch |
| 0xA1 | if icmplt | If two integers on stack are lt, branch |
| 0xA2 | if icmpge | If two integers on stack are ge, branch |
| 0xA3 | if icmpgt | If two integers on stack are gt, branch |
| 0xA4 | if icmple | If two integers on stack are le, branch |
| 0xA5 | if acmpeq | If two references on stack are the same, branch |
| 0xA6 | if acmpne | If two references on stack are different, branch |
| 0xA7 | goto | GOTO! |

Other Flow Control

| | | |
|-----------|-------------------|--|
| 0xA8 | jsr | Jump to subroutine (deprecated) |
| 0xA9 | ret | Return from subroutine (deprecated) |
| 0xAA | tableswitch | Branch using an indexed table of jump offsets |
| 0xAB | lookupswitch | Branch using a lookup-based table of jump offsets |
| 0xAC-0xB0 | [i,l,f,d,a]return | Return (int, long, float, double, reference) value |
| 0xB1 | return | Void return (exit method, return nothing) |
| 0xC6 | ifnull | If reference on stack is null |
| 0xC7 | ifnonnull | If reference on stack is not null |

Flow Control

aload 0

ldc 0

aaload

ldc "branch"

invokevirtual string, "equals",
[boolean, object]

ifne :branch

ldc "Not equal!"

aprintln

goto :end

label :branch

ldc "Equal!"

aprintln

label :end

returnvoid

| Depth | Value |
|-------|--------------------------------------|
| 0 | String[] {"branch"} |
| 1 | |
| 2 | |
| 3 | |
| 4 | |
| 5 | |

Flow Control

```
aload 0
ldc 0
aload
ldc "branch"
invokevirtual string, "equals",
    [boolean, object]

ifne :branch
ldc "Not equal!"
aprintln
goto :end
label :branch
ldc "Equal!"
aprintln
label :end
returnvoid
```

| Depth | Value |
|-------|--------------------|
| 0 | 0 |
| 1 | String[]{"branch"} |
| 2 | |
| 3 | |
| 4 | |
| 5 | |

Flow Control

```
aload 0
ldc 0
(aaload)
ldc "branch"
invokevirtual string, "equals",
    [boolean, object]

ifne :branch
ldc "Not equal!"
aprintln
goto :end
label :branch
ldc "Equal!"
aprintln
label :end
returnvoid
```

| Depth | Value |
|-------|-----------------|
| 0 | “branch” |
| 1 | |
| 2 | |
| 3 | |
| 4 | |
| 5 | |

Flow Control

```
aload 0
ldc 0
aload
ldc "branch"
invokevirtual string, "equals",
    [boolean, object]

ifne :branch
ldc "Not equal!"
aprintln
goto :end
label :branch
ldc "Equal!"
aprintln
label :end
returnvoid
```

| Depth | Value |
|-------|-----------------|
| 0 | “branch” |
| 1 | “branch” |
| 2 | |
| 3 | |
| 4 | |
| 5 | |

Flow Control

```
aload 0
ldc 0
aload
ldc "branch"
invokevirtual string, "equals",
    [boolean, object]
ifne :branch
ldc "Not equal!"
aprintln
goto :end
label :branch
ldc "Equal!"
aprintln
label :end
returnvoid
```

| Depth | Value |
|-------|-------|
| 0 | I |
| 1 | |
| 2 | |
| 3 | |
| 4 | |
| 5 | |


Flow Control

```
aload 0
ldc 0
aload
ldc "branch"
invokevirtual string, "equals",
    [boolean, object]
ifne :branch
ldc "Not equal!"
aprintln
goto :end
label :branch
ldc "Equal!"
aprintln
label :end
returnvoid
```

| Depth | Value |
|-------|-------|
| 0 | |
| 1 | |
| 2 | |
| 3 | |
| 4 | |
| 5 | |

Flow Control

```
aload 0
ldc 0
aload
ldc "branch"
invokevirtual string, "equals",
    [boolean, object]
ifne :branch
ldc "Not equal!"
aprintln
goto :end
label :branch
ldc "Equal!"
aprintln
label :end
returnvoid
```



| Depth | Value |
|-------|-------|
| 0 | |
| 1 | |
| 2 | |
| 3 | |
| 4 | |
| 5 | |

Flow Control

```
aload 0
ldc 0
aload
ldc "branch"
invokevirtual string, "equals",
    [boolean, object]

ifne :branch
ldc "Not equal!"
aprintln
goto :end
label :branch
ldc "Equal!"
aprintln
label :end
returnvoid
```

| Depth | Value |
|-------|-----------------|
| 0 | "Equal!" |
| 1 | |
| 2 | |
| 3 | |
| 4 | |
| 5 | |

Flow Control

```
aload 0
ldc 0
aload
ldc "branch"
invokevirtual string, "equals",
                        [boolean, object]

ifne :branch
ldc "Not equal!"
aprintln
goto :end
label :branch
ldc "Equal!"
aprintln
label :end
returnvoid
```

| Depth | Value |
|-------|-------|
| 0 | |
| 1 | |
| 2 | |
| 3 | |
| 4 | |
| 5 | |

Classes and Types

- Signatures!!!
 - Probably the most painful part
 - ...but not a big deal if you understand

Using Classes

| | | |
|------|-----------------|---|
| 0xB2 | getstatic | Fetch static field from class |
| 0xB3 | putstatic | Set static field in class |
| 0xB4 | getfield | Get instance field from object |
| 0xB5 | setfield | Set instance field in object |
| 0xB6 | invokevirtual | Invoke instance method on object |
| 0xB7 | invokespecial | Invoke constructor or “super” on object |
| 0xB8 | invokestatic | Invoke static method on class |
| 0xB9 | invokeinterface | Invoke interface method on object |
| 0xBA | invokedynamic | Invoke method dynamically on object (Java 7) |
| 0xBB | new | Construct new instance of object |
| 0xC0 | checkcast | Attempt to cast object to type |
| 0xC1 | instanceof | Push nonzero if object is instanceof specified type |

Using Classes

`new ArrayList`

```
dup
invokespecial ArrayList, '<init>',
               [void]
checkcast Collection
dup
ldc "first element"
invokeinterface Collection, 'add',
               [boolean, object]

pop
checkcast ArrayList
ldc 0
invokevirtual ArrayList, 'get',
               [object, int]

aprintln
returnvoid
```

| Depth | Value |
|-------|---|
| 0 | an ArrayList (uninitialized) |
| 1 | |
| 2 | |
| 3 | |
| 4 | |
| 5 | |

Using Classes

```
new ArrayList
dup
invokespecial ArrayList, '<init>',
    [void]
checkcast Collection
dup
ldc "first element"
invokeinterface Collection, 'add',
    [boolean, object]

pop
checkcast ArrayList
ldc 0
invokevirtual ArrayList, 'get',
    [object, int]

aprintln
returnvoid
```

| Depth | Value |
|-------|---|
| 0 | an ArrayList (uninitialized) |
| 1 | an ArrayList (uninitialized) |
| 2 | |
| 3 | |
| 4 | |
| 5 | |

Using Classes

```
new ArrayList
dup
invokespecial ArrayList, '<init>',
    [void]
checkcast Collection
dup
ldc "first element"
invokeinterface Collection, 'add',
    [boolean, object]

pop
checkcast ArrayList
ldc 0
invokevirtual ArrayList, 'get',
    [object, int]

println
returnvoid
```

| Depth | Value |
|-------|--------------|
| 0 | an ArrayList |
| 1 | |
| 2 | |
| 3 | |
| 4 | |
| 5 | |

Using Classes

```
new ArrayList
dup
invokespecial ArrayList, '<init>',
    [void]
(checkcast Collection)
dup
ldc "first element"
invokeinterface Collection, 'add',
    [boolean, object]

pop
checkcast ArrayList
ldc 0
invokevirtual ArrayList, 'get',
    [object, int]

aprintln
returnvoid
```

| Depth | Value |
|-------|---------------------|
| 0 | a Collection |
| 1 | |
| 2 | |
| 3 | |
| 4 | |
| 5 | |

Using Classes

```
new ArrayList
dup
invokespecial ArrayList, '<init>',
    [void]
checkcast Collection
dup
ldc "first element"
invokeinterface Collection, 'add',
    [boolean, object]

pop
checkcast ArrayList
ldc 0
invokevirtual ArrayList, 'get',
    [object, int]

aprintln
returnvoid
```

| Depth | Value |
|-------|---------------------|
| 0 | a Collection |
| 1 | a Collection |
| 2 | |
| 3 | |
| 4 | |
| 5 | |

Using Classes

```
new ArrayList
dup
invokespecial ArrayList, '<init>',
    [void]
checkcast Collection
dup
 ldc "first element"
 invokeinterface Collection, 'add',
    [boolean, object]

pop
checkcast ArrayList
 ldc 0
 invokevirtual ArrayList, 'get',
    [object, int]

println
returnvoid
```

| Depth | Value |
|-------|------------------------|
| 0 | “first element” |
| 1 | a Collection |
| 2 | a Collection |
| 3 | |
| 4 | |
| 5 | |

Using Classes

```
new ArrayList
dup
invokespecial ArrayList, '<init>',
                    [void]
checkcast Collection
dup
ldc "first element"
invokeinterface Collection, 'add',
                    [boolean, object]
pop
checkcast ArrayList
ldc 0
invokevirtual ArrayList, 'get',
                    [object, int]
println
returnvoid
```

| Depth | Value |
|-------|-----------------|
| 0 | I (true) |
| 1 | a Collection |
| 2 | |
| 3 | |
| 4 | |
| 5 | |

Using Classes

```
new ArrayList
dup
invokespecial ArrayList, '<init>',
                    [void]
checkcast Collection
dup
ldc "first element"
invokeinterface Collection, 'add',
                    [boolean, object]
(pop)
checkcast ArrayList
ldc 0
invokevirtual ArrayList, 'get',
                    [object, int]
println
returnvoid
```

| Depth | Value |
|-------|--------------|
| 0 | a Collection |
| 1 | |
| 2 | |
| 3 | |
| 4 | |
| 5 | |

Using Classes

```
new ArrayList
dup
invokespecial ArrayList, '<init>',
    [void]
checkcast Collection
dup
ldc "first element"
invokeinterface Collection, 'add',
    [boolean, object]

pop
(checkcast ArrayList)
ldc 0
invokevirtual ArrayList, 'get',
    [object, int]

aprintln
returnvoid
```

| Depth | Value |
|-------|---------------------|
| 0 | an ArrayList |
| 1 | |
| 2 | |
| 3 | |
| 4 | |
| 5 | |

Using Classes

```
new ArrayList
dup
invokespecial ArrayList, '<init>',
                    [void]
checkcast Collection
dup
ldc "first element"
invokeinterface Collection, 'add',
                    [boolean, object]

pop
checkcast ArrayList
(ldc 0)
invokevirtual ArrayList, 'get',
                    [object, int]

aprintln
returnvoid
```

| Depth | Value |
|-------|--------------|
| 0 | 0 |
| 1 | an ArrayList |
| 2 | |
| 3 | |
| 4 | |
| 5 | |

Using Classes

```
new ArrayList
dup
invokespecial ArrayList, '<init>',
                    [void]
checkcast Collection
dup
ldc "first element"
invokeinterface Collection, 'add',
                    [boolean, object]

pop
checkcast ArrayList
ldc 0
invokevirtual ArrayList, 'get',
                    [object, int]
println
returnvoid
```

| Depth | Value |
|-------|------------------------|
| 0 | “first element” |
| 1 | |
| 2 | |
| 3 | |
| 4 | |
| 5 | |

Using Classes

```
new ArrayList
dup
invokespecial ArrayList, '<init>',
                    [void]
checkcast Collection
dup
ldc "first element"
invokeinterface Collection, 'add',
                    [boolean, object]

pop
checkcast ArrayList
ldc 0
invokevirtual ArrayList, 'get',
                    [object, int]

(aprintln)
returnvoid
```

| Depth | Value |
|-------|-------|
| 0 | |
| 1 | |
| 2 | |
| 3 | |
| 4 | |
| 5 | |

Exceptions and Synchronization

| | | |
|------|--------------|--|
| - | trycatch | Table structure for a method indicating start/end of try/catch and logic to run on exception |
| 0xC2 | monitorenter | Enter synchronized block against object on stack |
| 0xC3 | monitorexit | Exit synchronized block (against same object) |

More Examples

- A simple loop
- Fibonacci

A Simple Loop

```
main do
  aload 0
  push_int 0
  aload
  label :top
  dup
  aprintln
  goto :top
  returnvoid
end
```

Fibonacci

```
public_static_method "fib", [], int, int do
  iload 0
  ldc 2
  if_icmpge :recurse
  iload 0
  ireturn
label :recurse
  iload 0
  ldc 1
  isub
  invokestatic this, "fib", [int, int]
  iload 0
  ldc 2
  isub
  invokestatic this, "fib", [int, int]
  iadd
  ireturn
end
```

Fibonacci

```
main do
  load_times
  istore 1

  ldc "Raw bytecode fib(45) performance:"
  aprintln

  label :top
  iload 1
  ifeq :done
  iinc 1, -1

  start_timing 2
  ldc 45
  invokestatic this, "fib", [int, int]
  pop
  end_timing 2

  ldc "Time: "
  aprintln
  lprintln 2
  goto :top

  label :done
  returnvoid
end
```

Fibonacci

```
main do
  load_times
  istore 1

  ldc "Raw bytecode fib(45) performance:"
  aprintln

  label :top
  iload 1
  ifeq :done
  iinc 1, -1

  start_timing 2
  ldc 45
  invokestatic this, "fib", [int, int]
  pop
  end_timing 2

  ldc "Time: "
  aprintln
  lprintln 2
  goto :top

  label :done
  returnvoid
end
```

Macros

Fibonacci

```
macro :load_times do
  aload 0
  ldc 0
  aload # number of times
  invokestatic JInteger, 'parseInt',
               [int, string]
end
```

Fibonacci

```
macro :start_timing do |i|  
  load_time  
  lstore i  
end
```


Fibonacci

```
macro :load_time do  
  invokestatic System, "currentTimeMillis", long  
end
```

Fibonacci

```
macro :end_timing do |i|  
  load_time  
  lload i  
  lsub  
  lstore i  
end
```

Fibonacci

```
macro :lprintln do |i|  
  getstatic System, "out", PrintStream  
  lload i  
  invokevirtual PrintStream, "println",  
    [void, long]  
end
```

ASM

- “All purpose bytecode manipulation and analysis framework.”
- De facto standard bytecode library
- <http://asm.ow2.org>

Basic Process

- Construct a ClassWriter
- Visit structure
 - Annotations, methods, fields, inner classes
- Write out bytes

Blah.java

```
public class Blah implements Cloneable {  
    private final String fieldName;  
  
    public Blah() {  
        fieldName = "hello";  
    }  
  
    public static Blah makeBlah() {  
        return new Blah();  
    }  
}
```

ClassWriter

```
ClassWriter cv = new ClassWriter(  
    ClassWriter.COMPUTE_MAXS |  
    ClassWriter.COMPUTE_FRAMES);
```

COMPUTE...what?

- COMPUTE_MAXS
 - ASM will calculate max stack/local vars
- COMPUTE_FRAMES
 - ASM will calculate Java 6 stack map
 - Hints to verifier that we've pre-validated stack contents (sort of)

Visit Class

```
cv.visit(  
    Opcodes.V1_6,  
    Opcodes.ACC_PUBLIC,  
    "Blah",  
    null,  
    "java/lang/Object",  
    new String[] {"java/lang/Cloneable"});
```

Opcodes

- Interface full of constants
 - Bytecodes
 - Visibility modifiers
 - Java versions
 - Other stuff

ACC_*

- Some you know
 - ACC_PUBLIC, ACC_ABSTRACT, etc
- Some you don't
 - ACC_BRIDGE, ACC_SYNTHETIC

Java Version

- VI_1 through VI_7
 - Sorry 1.0!

Class Names

"java/lang/Object"

```
packageClass.replaceAll('.', '/')
```

Visit Source

```
cv.visitSource(  
    "Blah.java",  
    "JSR-45 source map here");
```

Visit Annotation

```
AnnotationVisitor av = cv.visitAnnotation("some/Annotation", true);  
av.visitArray("name1", ...);  
av.visitEnum("name2", ...);  
av.visitEnd();
```

Blah.java

```
private final String fieldName;
```


Visit Field

```
FieldVisitor fv = cv.visitField(  
    Opcodes.ACC_PRIVATE | Opcodes.ACC_FINAL,  
    "fieldName",  
    "Ljava.lang.String;",  
    null);  
fv.visitAnnotation(...);  
fv.visitAttribute(...);  
fv.visitEnd();
```

Descriptor

"Ljava.lang.String;"
"(IF[Ljava.lang.Object;)V"

- Primitive types
 - B,C,S,I,J,F,D,Z,V
- Reference types
 - Lsome/Class;
- Array
 - Prefix with [

Blah.java

```
public Blah() {
```

```
    . . .
```

```
}
```

```
public static Blah makeBlah() {
```

```
    . . .
```

```
}
```

Visit Method

```
MethodVisitor construct = cv.visitMethod(  
    Opcodes.ACC_PUBLIC,  
    "<init>",  
    "()V",  
    null,  
    null);
```

```
MethodVisitor makeBlah = cv.visitMethod(  
    Opcodes.ACC_PUBLIC, ACC_STATIC,  
    "makeBlah",  
    "()LBlah;",  
    null,  
    null);
```

Special Methods

- <init>
 - Constructor
- <clinit>
 - Static initializer

MethodVisitor

- Visit annotation stuff
- Visit code
 - Bytecodes, frames, local vars, line nums
- Visit maxs
 - Pass bogus values if COMPUTE_MAXS

Blah.java

```
public Blah() {  
    fieldName = "hello";  
}
```

```
public static Blah makeBlah() {  
    return new Blah();  
}
```

Visit Method Body

```
construct.visitCode();  
construct.visitVarInsn(ALOAD, 0);  
construct.visitMethodInsn(INVOKEVIRTUAL,  
    "java/lang/Object",  
    "<init>",  
    "()V");  
construct.visitVarInsn(ALOAD, 0);  
construct.visitLdcInsn("hello");  
construct.visitFieldInsn(PUTFIELD,  
    "Blah",  
    "fieldName",  
    "Ljava/lang/String;");  
construct.visitInsn(RETURN);  
construct.visitMaxs(2, 1);  
construct.visitEnd();
```


ASMIifierClassVisitor

- Dump ASM visitor calls from .class file
- Very raw, but very useful

Blah.java

```
public class Blah implements Cloneable {  
    private final String fieldName;  
  
    public Blah() {  
        fieldName = "hello";  
    }  
  
    public static Blah makeBlah() {  
        return new Blah();  
    }  
}
```

```
~/oscon → java -cp asm-3.3.1.jar:asm-util-3.3.1.jar \
               org.objectweb.asm.util.ASMifierClassVisitor \
               Blah.class

import java.util.*;
import org.objectweb.asm.*;
import org.objectweb.asm.attrs.*;
public class BlahDump implements Opcodes {

    public static byte[] dump () throws Exception {

        ClassWriter cw = new ClassWriter(0);
        FieldVisitor fv;
        MethodVisitor mv;
        AnnotationVisitor av0;

        cw.visit(V1_6, ACC_PUBLIC + ACC_SUPER, "Blah", null, "java/lang/Object", new String[] { "java/lang/Cloneable" });

        {
            fv = cw.visitField(ACC_PRIVATE + ACC_FINAL, "fieldName", "Ljava/lang/String;", null, null);
            fv.visitEnd();
        }
        {
            mv = cw.visitMethod(ACC_PUBLIC, "<init>", "()V", null, null);
            mv.visitCode();
            mv.visitVarInsn(ALOAD, 0);
            mv.visitMethodInsn(INVOKESPECIAL, "java/lang/Object", "<init>", "()V");
            mv.visitVarInsn(ALOAD, 0);
            mv.visitLdcInsn("hello");
            mv.visitFieldInsn(PUTFIELD, "Blah", "fieldName", "Ljava/lang/String;");
            mv.visitInsn(RETURN);
            mv.visitMaxs(2, 1);
            mv.visitEnd();
        }
        {
            mv = cw.visitMethod(ACC_PUBLIC + ACC_STATIC, "makeBlah", "()LBlah;", null, null);
            mv.visitCode();
            mv.visitTypeInsn(NEW, "Blah");
            mv.visitInsn(DUP);
            mv.visitMethodInsn(INVOKESPECIAL, "Blah", "<init>", "()V");
            mv.visitInsn(ARETURN);
            mv.visitMaxs(2, 0);
            mv.visitEnd();
        }
        cw.visitEnd();

        return cw.toByteArray();
    }
}
```

Real-world Cases

- Reflection-free invocation
 - JRuby, Groovy, other languages
- Bytecoded data objects
 - Hibernate, other data layers
 - `java.lang.reflect.Proxy` and others
- Language compilers

Part Two:

JVM JIT

JIT

- Just-In-Time compilation
- Compiled when needed
 - Maybe immediately before execution
 - ...or when we decide it's important
 - ...or never?

Mixed-Mode

- Interpreted
 - Bytecode-walking
 - Artificial stack
- Compiled
 - Direct native operations
 - Native registers, memory, etc

Profiling

- Gather data about code while interpreting
 - Invariants (types, constants, nulls)
 - Statistics (branches, calls)
- Use that information to optimize
 - Educated guess?

Optimization

- Loop unrolling
- Lock coarsening
- Method inlining
- Dead code elimination
- Duplicate code elimination

The Golden Rule of Optimization

Don't do unnecessary work.

Perf Sinks

- Memory accesses
 - By far the biggest expense
- Calls
 - Opaque memory ref + branch
- Locks, volatile writes
 - Kills multi-cpu perf

Volatile?

- Each CPU maintains a memory cache
- Caches may be out of sync
 - If it doesn't matter, no problem
 - If it does matter, threads disagree!
- Volatile forces synchronization of cache
 - Across cores and to main memory

Inlining?

- Combine caller and callee into one unit
 - e.g. based on profile
 - Perhaps with a sanity check
- Optimize as a whole

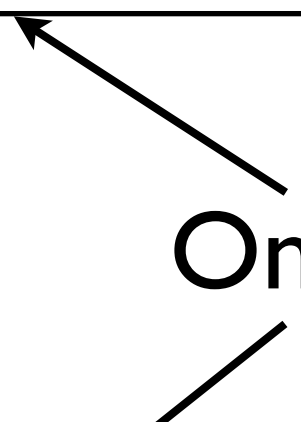
Inlining

```
int addAll(int max) {  
    int accum = 0;  
    for (int i = 0; i < max; i++) {  
        accum = add(accum, i);  
    }  
    return accum;  
}
```

```
int add(int a, int b) {  
    return a + b;  
}
```

Inlining

```
int addAll(int max) {  
    int accum = 0;  
    for (int i = 0; i < max; i++) {  
        accum = add(accum, i);  
    }  
    return accum;  
}
```



Only one target is ever seen

```
int add(int a, int b) {  
    return a + b;  
}
```

Inlining

```
int addAll(int max) {  
    int accum = 0;  
    for (int i = 0; i < max; i++) {  
        accum = accum + i;  
    }  
    return accum;  
}
```

Don't bother making a call



Call Site

- The place where you make a call
- Monomorphic (“one shape”)
 - Single target class
- Bimorphic (“two shapes”)
- Polymorphic (“many shapes”)
- Megamorphic (“you’re screwed”)

Blah.java

```
System.currentTimeMillis(); // static, monomorphic
```

```
List list1 = new ArrayList(); // constructor, monomorphic  
List list2 = new LinkedList();
```

```
for (List list : new List[]{ list1, list2 }) {  
    list.add("hello"); // bimorphic  
}
```

```
for (Object obj : new Object[]{ 'foo', list1, new Object() }) {  
    obj.toString(); // polymorphic  
}
```

Hotspot

- -client mode (C1) inlines, less aggressive
 - Fewer opportunities to optimize
- -server mode (C2) profiles, inlines
 - We'll focus on this
- Tiered mode combines them
 - -XX:+TieredCompilation

Hotspot Inlining

- Profile to find “hot spots”
 - Largely focused around call sites
 - Profile until 10k calls
- Inline mono/bimorphic calls
- Other mechanisms for polymorphic calls

Now it gets fun!

Monitoring the JIT

- Dozens of flags
 - PrintCompilation
 - PrintInlining
 - LogCompilation
 - PrintAssembly
- Some in product, some in debug...

```
public class Accumulator {  
    public static void main(String[] args) {  
        int max = Integer.parseInt(args[0]);  
        System.out.println(addAll(max));  
    }  
  
    static int addAll(int max) {  
        int accum = 0;  
        for (int i = 0; i < max; i++) {  
            accum = add(accum, i);  
        }  
        return accum;  
    }  
  
    static int add(int a, int b) {  
        return a + b;  
    }  
}
```

```
~/oscon → java -version  
openjdk version "1.7.0-internal"  
OpenJDK Runtime Environment (build 1.7.0-internal-b00)  
OpenJDK 64-Bit Server VM (build 21.0-b17, mixed mode)
```

```
~/oscon → javac Accumulator.java
```

```
~/oscon → java Accumulator 1000  
499500
```


PrintCompilation


- -XX:+PrintCompilation
- Print methods as they are jitted
 - Class + name + size

```
~/oscon → java -XX:+PrintCompilation Accumulator 1000
 1      java.lang.String::hashCode (64 bytes)
 2      java.math.BigInteger::mulAdd (81 bytes)
 3      java.math.BigInteger::multiplyToLen (219 bytes)
 4      java.math.BigInteger::addOne (77 bytes)
 5      java.math.BigInteger::squareToLen (172 bytes)
 6      java.math.BigInteger::primitiveLeftShift (79 bytes)
 7      java.math.BigInteger::montReduce (99 bytes)
 8      sun.security.provider.SHA::implCompress (491 bytes)
 9      java.lang.String::charAt (33 bytes)
499500
```

```
~/oscon → java -XX:+PrintCompilation Accumulator 1000
1      java.lang.String::hashCode (64 bytes)
2      java.math.BigInteger::mulAdd (81 bytes)
3      java.math.BigInteger::multiplyToLen (219 bytes)
4      java.math.BigInteger::addOne (77 bytes)
5      java.math.BigInteger::squareToLen (172 bytes)
6      java.math.BigInteger::primitiveLeftShift (79 bytes)
7      java.math.BigInteger::montReduce (99 bytes)
8      sun.security.provider.SHA::implCompress (491 bytes)
9      java.lang.String::charAt (33 bytes)
499500
```

Where's our methods?!


```
~/oscon → java -XX:+PrintCompilation Accumulator 1000
1      java.lang.String::hashCode (64 bytes)
2      java.math.BigInteger::mulAdd (81 bytes)
3      java.math.BigInteger::multiplyToLen (219 bytes)
4      java.math.BigInteger::addOne (77 bytes)
5      java.math.BigInteger::squareToLen (172 bytes)
6      java.math.BigInteger::primitiveLeftShift (79 bytes)
7      java.math.BigInteger::montReduce (99 bytes)
8      sun.security.provider.SHA::implCompress (491 bytes)
9      java.lang.String::charAt (33 bytes)
499500
```



Where's our methods?!

...remember... 10k calls

10k loop, 10k calls to add



```
~/oscon → java -XX:+PrintCompilation Accumulator 10000
1      java.lang.String::hashCode (64 bytes)
2      java.math.BigInteger::mulAdd (81 bytes)
3      java.math.BigInteger::multiplyToLen (219 bytes)
4      java.math.BigInteger::addOne (77 bytes)
5      java.math.BigInteger::squareToLen (172 bytes)
6      java.math.BigInteger::primitiveLeftShift (79 bytes)
7      java.math.BigInteger::montReduce (99 bytes)
8      sun.security.provider.SHA::implCompress (491 bytes)
9      java.lang.String::charAt (33 bytes)
10     Accumulator::add (4 bytes)
49995000
```



Hooray!

What's this stuff?

```
1 java.lang.String::hashCode (64 bytes)
2 java.math.BigInteger::mulAdd (81 bytes)
3 java.math.BigInteger::multiplyToLen (219 bytes)
4 java.math.BigInteger::addOne (77 bytes)
5 java.math.BigInteger::squareToLen (172 bytes)
6 java.math.BigInteger::primitiveLeftShift (79 bytes)
7 java.math.BigInteger::montReduce (99 bytes)
8 sun.security.provider.SHA::implCompress (491 bytes)
9 java.lang.String::charAt (33 bytes)
```

What's this stuff?

```
1  java.lang.String::hashCode (64 bytes)
2  java.math.BigInteger::mulAdd (81 bytes)
3  java.math.BigInteger::multiplyToLen (219 bytes)
4  java.math.BigInteger::addOne (77 bytes)
5  java.math.BigInteger::squareToLen (172 bytes)
6  java.math.BigInteger::primitiveLeftShift (79 bytes)
7  java.math.BigInteger::montReduce (99 bytes)
8  sun.security.provider.SHA::implCompress (491 bytes)
9  java.lang.String::charAt (33 bytes)
```

Class loading, security, other boot logic.

What if you see this...

```
~/oscon → java -client -XX:+PrintCompilation Accumulator 1000
 1      java.lang.String::hashCode (64 bytes)
 2      java.math.BigInteger::mulAdd (81 bytes)
 3      java.math.BigInteger::multiplyToLen (219 bytes)
 4      java.math.BigInteger::addOne (77 bytes)
 5      java.math.BigInteger::squareToLen (172 bytes)
 5 made not entrant java.math.BigInteger::squareToLen (172 bytes)
 7      java.math.BigInteger::montReduce (99 bytes)
 6      java.math.BigInteger::primitiveLeftShift (79 bytes)
 8      java.math.BigInteger::squareToLen (172 bytes)
 9      sun.security.provider.SHA::implCompress (491 bytes)
10      java.lang.String::charAt (33 bytes)
499500
```


Optimistic Compiler

- Assume profile is accurate
- Aggressively optimize based on profile
- Bail out if we're wrong
 - And hope that we're usually right

Deoptimization


- Bail out of running code
- Monitoring flags describe process
 - “uncommon trap” - we were wrong
 - “not entrant” - don’t let anyone enter
 - “zombie” - on its way to deadness

What if you see this...

```
20      java.math.BigInteger::addOne (77 bytes)
21      java.math.BigInteger::squareToLen (172 bytes)
22      java.math.BigInteger::primitiveLeftShift (79 bytes)
---    n  java.lang.System::arraycopy (static)
24      sun.security.provider.SHA::implCompress (491 bytes)
23      java.math.BigInteger::montReduce (99 bytes)
25      java.lang.String$CaseInsensitiveComparator::compare (115
bytes)
26      java.lang.Character::toLowerCase (162 bytes)
```

What if you see this...

```
20      java.math.BigInteger::addOne (77 bytes)
21      java.math.BigInteger::squareToLen (172 bytes)
22      java.math.BigInteger::primitiveLeftShift (79 bytes)
---      n      java.lang.System::arraycopy (static)
24      sun.security.provider.SHA::implCompress (491 bytes)
23      java.math.BigInteger::montReduce (99 bytes)
25      java.lang.String$CaseInsensitiveComparator::compare (115
bytes)
26      java.lang.Character::toLowerCase (162 bytes)
```



Native calls don't compile, may be
intrinsic. We'll come back to that.

PrintInlining

- `-XX:+UnlockDiagnosticVMOptions`
`-XX:+PrintInlining`
- Display hierarchy of inlined methods
- Include reasons for not inlining
- More, better output on OpenJDK 7

```
~/oscon → java -XX:+UnlockDiagnosticVMOptions \  
>          -XX:+PrintInlining \  
>          Accumulator 10000  
49995000
```

```
~/oscon → java -XX:+UnlockDiagnosticVMOptions \  
>          -XX:+PrintInlining \  
>          Accumulator 10000  
49995000
```

Um...I don't see anything inlining...

```
public class Accumulator {
    public static void main(String[] args) {
        int max = Integer.parseInt(args[0]);
        System.out.println(addAll(max));
    }

    static int addAll(int max) {
        int accum = 0;
        for (int i = 0; i < max; i++) {
            accum = add(accum, i);
        }
        return accum;
    }

    static int add(int a, int b) {
        return a + b;
    }
}
```

Called once

Called 10k times...

...but no calls to inline!


```
public class Accumulator2 {  
    public static void main(String[] args) {  
        int max = Integer.parseInt(args[0]);  
        System.out.println(  
            new Accumulator2().addAllSqrts(max));  
    }  
  
    double addAllSqrts(int max) {  
        double accum = 0;  
        for (int i = 0; i < max; i++) {  
            accum = addSqrt(accum, i);  
        }  
        return accum;  
    }  
  
    double addSqrt(double a, int b) {  
        return a + sqrt(b);  
    }  
  
    double sqrt(int b) {  
        return Math.sqrt(b);  
    }  
}
```


```
~/oscon → java -XX:+UnlockDiagnosticVMOptions \  
> -XX:+PrintCompilation \  
> -XX:+PrintInlining \  
> Accumulator2 10000
```

```
...  
89      2      Accumulator2::addSqrt (8 bytes)  
           @ 3    Accumulator2::sqrt (6 bytes)   inline (hot)  
           @ 2    java.lang.Math::sqrt (5 bytes) (intrinsic)  
89      3      Accumulator2::sqrt (6 bytes)  
           @ 2    java.lang.Math::sqrt (5 bytes) (intrinsic)  
666616.4591971082
```

A hot spot!



Calls treated specially by JIT



Intrinsic?

- Known to the JIT
 - Don't inline the bytecode
 - Do perform a specific native operation
 - e.g. kernel-level memory operation
 - e.g. optimized sqrt in machine code

Did Someone Say
MACHINE CODE?

The Red Pill

- Knowing code compiles is good
- Knowing code inlines is better
- Seeing the actual assembly is best!

Caveat

- I don't really know assembly.
- But I fake it really well.

PrintAssembly

- `-XX:+PrintAssembly`
- Google “hotspot printassembly”
- <http://wikis.sun.com/display/HotSpotInternals/PrintAssembly>
- Assembly dumping plugins

```
~/oscon → java -XX:+UnlockDiagnosticVMOptions \  
>           -XX:+PrintAssembly \  
>           Accumulator 10000  
OpenJDK 64-Bit Server VM warning: PrintAssembly is enabled;  
turning on DebugNonSafepoints to gain additional output  
Loaded disassembler from hsdis-amd64.dylib  
...
```


Decoding compiled method 11343cbd0:

Code:

[Disassembling for mach='i386:x86-64']

[Entry Point]

[Verified Entry Point]

[Constants]

{method} 'add' '(II)I' in 'Accumulator'

parm0: rsi = int

parm1: rdx = int

[sp+0x20] (sp of caller)

11343cd00: push %rbp

11343cd01: sub \$0x10,%rsp

11343cd05: nop

; *synchronization entry

; - Accumulator::add@-1 (line 16)

11343cd06: mov %esi,%eax

11343cd08: add %edx,%eax

; *iadd

; - Accumulator::add@2 (line 16)

11343cd0a: add \$0x10,%rsp

11343cd0e: pop %rbp

11343cd0f: test %eax,-0x1303fd15(%rip) # 1003fd000

; {poll_return}

11343cd15: retq

Woah there, buddy.

x86_64 Assembly I01

| | |
|-----------------------|--------------------------------|
| add | Two's complement add |
| sub | ...subtract |
| mov* | Move data from a to b |
| jmp | goto |
| je, jne, jl, jge, ... | Jump if ==, !=, <, >=, ... |
| push, pop | Push/pop to/from call stack |
| call*, ret* | Call or return from subroutine |
| eax, ebx, esi, ... | Registers |
| rdx, rbx, rsi, ... | 64-bit registers |

Register Machine

- Instead of stack moves, we have “slots”
- Move data into slots
- Call operations that work with slots
- Get new data out of slots
- JVM stack, locals end up as register ops

Stack?

- Native code has a stack too
 - Maintains context from call to call
- Calling conventions
 - Caller preserves registers?
 - Callee preserves registers?
 - Many different styles

| | |
|---|---------------------------------------|
| Decoding compiled method 11343cbd0: | <= address of new compiled code |
| Code: | |
| [Disassembling for mach='i386:x86-64'] | <= architecture |
| [Entry Point] | |
| [Verified Entry Point] | |
| [Constants] | |
| # {method} 'add' '(II)I' in 'Accumulator' | <= method, signature, class |
| # parm0: rsi = int | <= first parm to method goes in rsi |
| # parm1: rdx = int | <= second parm goes in rdx |
| # [sp+0x20] (sp of caller) | <= caller's pointer into native stack |

| | |
|---|-----------------------------------|
| 11343cd00: push %rbp | |
| 11343cd01: sub \$0x10,%rsp | |
| 11343cd05: nop | ; *synchronization entry |
| | ; - Accumulator::add@-1 (line 16) |
| 11343cd06: mov %esi,%eax | |
| 11343cd08: add %edx,%eax | ; *iadd |
| | ; - Accumulator::add@2 (line 16) |
| 11343cd0a: add \$0x10,%rsp | |
| 11343cd0e: pop %rbp | |
| 11343cd0f: test %eax,-0x1303fd15(%rip) | # 1003fd000 |
| | ; {poll_return} |
| 11343cd15: retq | |

rbp points at current stack frame, so we save it off.

| | | |
|-----------------|------------------------|-----------------------------------|
| 11343cd00: push | %rbp | |
| 11343cd01: sub | \$0x10,%rsp | |
| 11343cd05: nop | | ; *synchronization entry |
| | | ; - Accumulator::add@-1 (line 16) |
| 11343cd06: mov | %esi,%eax | |
| 11343cd08: add | %edx,%eax | ; *iadd |
| | | ; - Accumulator::add@2 (line 16) |
| 11343cd0a: add | \$0x10,%rsp | |
| 11343cd0e: pop | %rbp | |
| 11343cd0f: test | %eax,-0x1303fd15(%rip) | # 1003fd000 |
| | | ; {poll_return} |
| 11343cd15: retq | | |

Two args, so we bump stack pointer by 0x10.

| | | |
|-----------------------|------------------------|-----------------------------------|
| 11343cd00: push | %rbp | |
| 11343cd01: sub | \$0x10,%rsp | |
| 11343cd05: <u>nop</u> | | ; *synchronization entry |
| | | ; - Accumulator::add@-1 (line 16) |
| 11343cd06: mov | %esi,%eax | |
| 11343cd08: add | %edx,%eax | ; *iadd |
| | | ; - Accumulator::add@2 (line 16) |
| 11343cd0a: add | \$0x10,%rsp | |
| 11343cd0e: pop | %rbp | |
| 11343cd0f: test | %eax,-0x1303fd15(%rip) | # 1003fd000 |
| | | ; {poll_return} |
| 11343cd15: retq | | |

Do nothing, e.g. to memory-align code.

| | | |
|-----------------|------------------------|-----------------------------------|
| 11343cd00: push | %rbp | |
| 11343cd01: sub | \$0x10,%rsp | |
| 11343cd05: nop | | ;*synchronization entry |
| | | ; - Accumulator::add@-1 (line 16) |
| 11343cd06: mov | %esi,%eax | |
| 11343cd08: add | %edx,%eax | ;*iadd |
| | | ; - Accumulator::add@2 (line 16) |
| 11343cd0a: add | \$0x10,%rsp | |
| 11343cd0e: pop | %rbp | |
| 11343cd0f: test | %eax,-0x1303fd15(%rip) | # 1003fd000 |
| | | ; {poll_return} |
| 11343cd15: retq | | |

At the “-l” instruction of our add() method...
i.e. here we go!

| | | |
|-----------------|------------------------|-----------------------------------|
| 11343cd00: push | %rbp | |
| 11343cd01: sub | \$0x10,%rsp | |
| 11343cd05: nop | | ;*synchronization entry |
| | | ; - Accumulator::add@-1 (line 16) |
| 11343cd06: mov | %esi,%eax | |
| 11343cd08: add | %edx,%eax | ;*iadd |
| | | ; - Accumulator::add@2 (line 16) |
| 11343cd0a: add | \$0x10,%rsp | |
| 11343cd0e: pop | %rbp | |
| 11343cd0f: test | %eax,-0x1303fd15(%rip) | # 1003fd000 |
| | | ; {poll_return} |
| 11343cd15: retq | | |

Move parm1 into eax.

| | | |
|-----------------|------------------------|-----------------------------------|
| 11343cd00: push | %rbp | |
| 11343cd01: sub | \$0x10,%rsp | |
| 11343cd05: nop | | ;*synchronization entry |
| | | ; - Accumulator::add@-1 (line 16) |
| 11343cd06: mov | %esi,%eax | |
| 11343cd08: add | %edx,%eax | ;*iadd |
| | | ; - Accumulator::add@2 (line 16) |
| 11343cd0a: add | \$0x10,%rsp | |
| 11343cd0e: pop | %rbp | |
| 11343cd0f: test | %eax,-0x1303fd15(%rip) | # 1003fd000 |
| | | ; {poll_return} |
| 11343cd15: retq | | |

Add parm0 and parm1, store result in eax.

| | | |
|-----------------|------------------------|-----------------------------------|
| 11343cd00: push | %rbp | |
| 11343cd01: sub | \$0x10,%rsp | |
| 11343cd05: nop | | ; *synchronization entry |
| | | ; - Accumulator::add@-1 (line 16) |
| 11343cd06: mov | %esi,%eax | |
| 11343cd08: add | %edx,%eax | (;*iadd) |
| | | ; - Accumulator::add@2 (line 16) |
| 11343cd0a: add | \$0x10,%rsp | |
| 11343cd0e: pop | %rbp | |
| 11343cd0f: test | %eax,-0x1303fd15(%rip) | # 1003fd000 |
| | | ; {poll_return} |
| 11343cd15: retq | | |

How nice, Hotspot shows us this is our “iadd” op!

| | | |
|-----------------|------------------------|-----------------------------------|
| 11343cd00: push | %rbp | |
| 11343cd01: sub | \$0x10,%rsp | |
| 11343cd05: nop | | ;*synchronization entry |
| | | ; - Accumulator::add@-1 (line 16) |
| 11343cd06: mov | %esi,%eax | |
| 11343cd08: add | %edx,%eax | ;*iadd |
| | | ; - Accumulator::add@2 (line 16) |
| 11343cd0a: add | \$0x10,%rsp | |
| 11343cd0e: pop | %rbp | |
| 11343cd0f: test | %eax,-0x1303fd15(%rip) | # 1003fd000 |
| | | ; {poll_return} |
| 11343cd15: retq | | |

Put stack pointer back where it was.

| | | |
|-----------------|------------------------|-----------------------------------|
| 11343cd00: push | %rbp | |
| 11343cd01: sub | \$0x10,%rsp | |
| 11343cd05: nop | | ; *synchronization entry |
| | | ; - Accumulator::add@-1 (line 16) |
| 11343cd06: mov | %esi,%eax | |
| 11343cd08: add | %edx,%eax | ; *iadd |
| | | ; - Accumulator::add@2 (line 16) |
| 11343cd0a: add | \$0x10,%rsp | |
| 11343cd0e: pop | %rbp | |
| 11343cd0f: test | %eax,-0x1303fd15(%rip) | # 1003fd000 |
| | | ; {poll_return} |
| 11343cd15: retq | | |

Restore rbp from stack.

| | | |
|-----------------|------------------------|-----------------------------------|
| 11343cd00: push | %rbp | |
| 11343cd01: sub | \$0x10,%rsp | |
| 11343cd05: nop | | ;*synchronization entry |
| | | ; - Accumulator::add@-1 (line 16) |
| 11343cd06: mov | %esi,%eax | |
| 11343cd08: add | %edx,%eax | ;*iadd |
| | | ; - Accumulator::add@2 (line 16) |
| 11343cd0a: add | \$0x10,%rsp | |
| 11343cd0e: pop | %rbp | |
| 11343cd0f: test | %eax,-0x1303fd15(%rip) | # 1003fd000 |
| | | ; {poll_return} |
| 11343cd15: retq | | |

Poll a “safepoint”...give JVM a chance to GC, etc.

| | | |
|-----------------|------------------------|-----------------------------------|
| 11343cd00: push | %rbp | |
| 11343cd01: sub | \$0x10,%rsp | |
| 11343cd05: nop | | ; *synchronization entry |
| | | ; - Accumulator::add@-1 (line 16) |
| 11343cd06: mov | %esi,%eax | |
| 11343cd08: add | %edx,%eax | ; *iadd |
| | | ; - Accumulator::add@2 (line 16) |
| 11343cd0a: add | \$0x10,%rsp | |
| 11343cd0e: pop | %rbp | |
| 11343cd0f: test | %eax,-0x1303fd15(%rip) | # 1003fd000 |
| | | ; {poll_return} |
| 11343cd15: | retq | |

All done!

Things to Watch For

- CALL operations
 - Indicate something failed to inline
- LOCK operations
 - Cache-busting, e.g. volatility

CALL

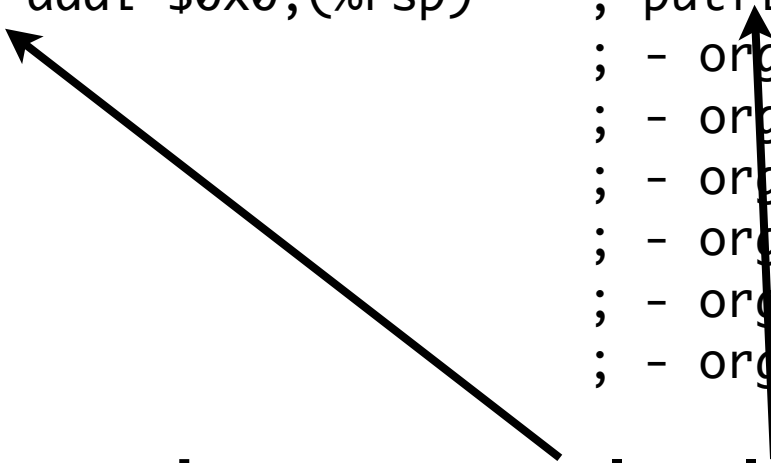
```
1134858f5: xchg    %ax,%ax
1134858f7: callq   113414aa0    ; OopMap{off=316}
                          ; *invokespecial addAsBignum
                          ; - org.jruby.RubyFixnum::addFixnum@29 (line 348)
                          ; {optimized virtual_call}
1134858fc: jmpq    11348586d
```

Ruby integer adds might overflow into Bignum, leading to addAsBignum call. In this case, it's never called, so Hotspot emits callq assuming we won't hit it.

LOCK

Code from a RubyBasicObject's default constructor.

```
11345d823: mov    0x70(%r8),%r9d    ;*getstatic NULL_OBJECT_ARRAY
                                ; - org.jruby.RubyBasicObject::<init>@5 (line 76)
                                ; - org.jruby.RubyObject::<init>@2 (line 118)
                                ; - org.jruby.RubyNumeric::<init>@2 (line 111)
                                ; - org.jruby.RubyInteger::<init>@2 (line 95)
                                ; - org.jruby.RubyFixnum::<init>@5 (line 112)
                                ; - org.jruby.RubyFixnum::newFixnum@25 (line 173)
11345d827: mov    %r9d,0x14(%rax)
11345d82b: lock addl $0x0,(%rsp)    ;*putfield varTable
                                ; - org.jruby.RubyBasicObject::<init>@8 (line 76)
                                ; - org.jruby.RubyObject::<init>@2 (line 118)
                                ; - org.jruby.RubyNumeric::<init>@2 (line 111)
                                ; - org.jruby.RubyInteger::<init>@2 (line 95)
                                ; - org.jruby.RubyFixnum::<init>@5 (line 112)
                                ; - org.jruby.RubyFixnum::newFixnum@25 (line 173)
```



Why are we doing a volatile write in the constructor?

LOCK

```
public class RubyBasicObject ... {  
    private static final boolean DEBUG = false;  
    private static final Object[] NULL_OBJECT_ARRAY = new Object[0];  
  
    // The class of this object  
    protected transient RubyClass metaClass;  
  
    // zeroed by jvm  
    protected int flags;  
  
    // variable table, lazily allocated as needed (if needed)  
    private volatile Object[] varTable = NULL_OBJECT_ARRAY;
```

Maybe it's not such a good idea to pre-init a volatile?



LOCK

```
~/projects/jruby → git log 2f935de1e40bfd8b29b3a74eaed699e519571046 -1 | cat  
commit 2f935de1e40bfd8b29b3a74eaed699e519571046  
Author: Charles Oliver Nutter <headius@headius.com>  
Date: Tue Jun 14 02:59:41 2011 -0500
```

Do not eagerly initialize volatile varTable field in RubyBasicObject;
speeds object creation significantly.

LEVEL UP!

What have we learned?

- How to emit bytecode
- How to read bytecode
- How bytecode execution works
- How to monitor the Hotspot JIT
- How to find problems from asm code

You're no dummy now.
;-)

Thank you!

- headius@headius.com, @headius
- http://blog.headius.com
- http://github.com/headius/bitescript
- “java virtual machine specification”
- “jvm opcodes”